

Practical Machine Learning - Prediction Assignment Writeup

1. SUMMARY

Devices such as Jawbone Up, Nike FuelBand, and Fitbit make possible to collect a large amount of data about personal activity relatively inexpensively.

One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it.

The provided dataset contained data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants. They were asked to perform barbell lifts correctly and incorrectly, in 5 different ways.

- Training data source : <https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>
- Test data source : <https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>
- Project source : <http://groupware.les.inf.puc-rio.br/har>

In this assignment, the goal was to build a model to accurately predict the manner in which people did their exercises. To accomplish this task, it was used random forests method and the **bigrf** package.

The overall prediction accuracy was approximately **99.3%**, which was a good and encouraging result. As reference, the accuracy achieved by in the original paper which used the proposed data set was **98.2%**. The out of sample error in the prediction was low, in the order of **0.726%**.

2. DOWNLOADING THE DATA AND LOADING INTO R

```
setwd("/Volumes/Documentos importantes/Coursera/8 - Practical Machine Learning/Quizzes and Project")
download.file("https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv",
  destfile = "pml-training.csv", method = "curl")
download.file("https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv",
  destfile = "pml-testing.csv", method = "curl")
train <- read.csv("pml-training.csv", header = T, na.strings = c("NA", "#DIV/0!",
  ""))
test <- read.csv("pml-testing.csv", header = T, na.strings = c("NA", "#DIV/0!",
  ""))
```

3. PARTITIONING THE TRAINING DATA SET

The training data set (object *train*) was partitioned in two data sets :

- 60% for algorithm development (*train2* object).
- 40% for algorithm testing before application to test dataset and out of sample error estimation(*test2* object).

```
if("caret" %in% rownames(installed.packages()) == FALSE) {install.packages("caret")}
library(caret)
```

```
## Loading required package: lattice
## Loading required package: ggplot2
```

```
set.seed(1)
split <- createDataPartition(y=train$classe, p=0.6, list=FALSE)
train2 <- train[split, ]; test2 <- train[-split, ]
dim(train2); dim(test2)
```

```
## [1] 11776 160
```

```
## [1] 7846 160
```

4. CLEANING AND CHECKING THE DATA SETS

The classification of the types of exercises should be performed based on the right variables, such as data from sensors on the belt, forearm, arm, and dumbbell. Additionally, the result of the prediction should accurately classify the types of exercises as **correctly performed** (class A) or **performed incorrectly** (classes B, C, D, and E).

The mentioned variables follow the presented name patterns :

- gyros_XXX_x, gyros_XXX_y, gyros_XXX_z
- accel_XXX_x, accel_XXX_y, accel_XXX_z
- total_accel_XXX
- magnet_XXX_x, magnet_XXX_y, magnet_XXX_z
- roll_XXX
- pitch_XXX
- yaw_XXX

So, the predictions were extrated using :

```
preds <- c(grep("^accel", names(train)), grep("^gyros", names(train)), grep("^magnet",
  names(train)), grep("^roll", names(train)), grep("^pitch", names(train)),
  grep("^yaw", names(train)), grep("^total", names(train)), grep("classe",
    names(train)))
train2_pred <- train2[, preds]
test2_pred <- test2[, preds]
```

It was possible to confirm that the remaining predictors were actually suitable, using the `nearZeroVar` function from `Caret` package. No remaining variable has variance close to zero :

```
nzv_train2<-nearZeroVar(train2_pred,saveMetrics=T)
nzv_test2<-nearZeroVar(test2_pred,saveMetrics=T)
nzv_train2; nzv_test2
```

##	freqRatio	percentUnique	zeroVar	nzv
## accel_belt_x	1.000000	1.34171196	FALSE	FALSE
## accel_belt_y	1.125140	1.16338315	FALSE	FALSE
## accel_belt_z	1.088561	2.39470109	FALSE	FALSE
## accel_arm_x	1.097087	6.43682065	FALSE	FALSE
## accel_arm_y	1.185484	4.40726902	FALSE	FALSE
## accel_arm_z	1.025974	6.41134511	FALSE	FALSE
## accel_dumbbell_x	1.045918	3.48165761	FALSE	FALSE
## accel_dumbbell_y	1.040541	3.87228261	FALSE	FALSE
## accel_dumbbell_z	1.129252	3.30332880	FALSE	FALSE
## accel_forearm_x	1.140351	6.60665761	FALSE	FALSE
## accel_forearm_y	1.032258	8.23709239	FALSE	FALSE
## accel_forearm_z	1.087912	4.68750000	FALSE	FALSE
## gyros_belt_x	1.058824	1.05298913	FALSE	FALSE
## gyros_belt_y	1.148268	0.55197011	FALSE	FALSE
## gyros_belt_z	1.126091	1.32472826	FALSE	FALSE
## gyros_arm_x	1.000000	5.27343750	FALSE	FALSE
## gyros_arm_y	1.376147	3.04008152	FALSE	FALSE
## gyros_arm_z	1.023739	1.98709239	FALSE	FALSE
## gyros_dumbbell_x	1.014045	1.90217391	FALSE	FALSE
## gyros_dumbbell_y	1.247887	2.22486413	FALSE	FALSE
## gyros_dumbbell_z	1.099448	1.63043478	FALSE	FALSE
## gyros_forearm_x	1.085443	2.35224185	FALSE	FALSE
## gyros_forearm_y	1.109091	6.02921196	FALSE	FALSE
## gyros_forearm_z	1.081633	2.44565217	FALSE	FALSE

## magnet_belt_x	1.046512	2.50509511	FALSE	FALSE
## magnet_belt_y	1.116711	2.36922554	FALSE	FALSE
## magnet_belt_z	1.021505	3.60054348	FALSE	FALSE
## magnet_arm_x	1.039216	11.02241848	FALSE	FALSE
## magnet_arm_y	1.132075	7.19259511	FALSE	FALSE
## magnet_arm_z	1.000000	10.59782609	FALSE	FALSE
## magnet_dumbbell_x	1.065421	8.82302989	FALSE	FALSE
## magnet_dumbbell_y	1.326531	6.85292120	FALSE	FALSE
## magnet_dumbbell_z	1.096491	5.62160326	FALSE	FALSE
## magnet_forearm_x	1.063830	11.99048913	FALSE	FALSE
## magnet_forearm_y	1.075472	15.26834239	FALSE	FALSE
## magnet_forearm_z	1.026316	13.40013587	FALSE	FALSE
## roll_belt	1.043557	8.74660326	FALSE	FALSE
## roll_arm	53.179487	19.24252717	FALSE	FALSE
## roll_dumbbell	1.075000	87.48301630	FALSE	FALSE
## roll_forearm	12.729730	14.97112772	FALSE	FALSE
## pitch_belt	1.026786	13.79076087	FALSE	FALSE
## pitch_arm	74.071429	22.46093750	FALSE	FALSE
## pitch_dumbbell	2.425000	85.44497283	FALSE	FALSE
## pitch_forearm	61.973684	20.88145380	FALSE	FALSE
## yaw_belt	1.104027	14.55502717	FALSE	FALSE
## yaw_arm	32.920635	21.22961957	FALSE	FALSE
## yaw_dumbbell	1.052632	87.02445652	FALSE	FALSE
## yaw_forearm	15.798658	14.17289402	FALSE	FALSE
## total_accel_belt	1.044661	0.23777174	FALSE	FALSE
## total_accel_arm	1.092453	0.55197011	FALSE	FALSE
## total_accel_dumbbell	1.085185	0.35665761	FALSE	FALSE
## total_accel_forearm	1.151194	0.58593750	FALSE	FALSE
## classe	1.469065	0.04245924	FALSE	FALSE

##	freqRatio	percentUnique	zeroVar	nzv
## accel_belt_x	1.057508	1.93729289	FALSE	FALSE
## accel_belt_y	1.097638	1.64414989	FALSE	FALSE
## accel_belt_z	1.062874	3.46673464	FALSE	FALSE
## accel_arm_x	1.083333	9.29135865	FALSE	FALSE
## accel_arm_y	1.077778	6.44914606	FALSE	FALSE
## accel_arm_z	1.265306	9.22763191	FALSE	FALSE
## accel_dumbbell_x	1.014815	4.66479735	FALSE	FALSE
## accel_dumbbell_y	1.072165	5.48049962	FALSE	FALSE
## accel_dumbbell_z	1.139785	5.02166709	FALSE	FALSE
## accel_forearm_x	1.025000	9.60999235	FALSE	FALSE
## accel_forearm_y	1.097561	11.90415498	FALSE	FALSE
## accel_forearm_z	1.093750	6.66581698	FALSE	FALSE
## gyros_belt_x	1.058380	1.45296967	FALSE	FALSE
## gyros_belt_y	1.137553	0.76472088	FALSE	FALSE
## gyros_belt_z	1.017981	1.95003824	FALSE	FALSE
## gyros_arm_x	1.040201	7.68544481	FALSE	FALSE
## gyros_arm_y	1.590426	4.46087178	FALSE	FALSE
## gyros_arm_z	1.267380	2.68926842	FALSE	FALSE
## gyros_dumbbell_x	1.047809	2.80397655	FALSE	FALSE
## gyros_dumbbell_y	1.232365	3.16084629	FALSE	FALSE
## gyros_dumbbell_z	1.000000	2.30690798	FALSE	FALSE
## gyros_forearm_x	1.019324	3.31379047	FALSE	FALSE
## gyros_forearm_y	1.065359	8.66683660	FALSE	FALSE
## gyros_forearm_z	1.188172	3.19908233	FALSE	FALSE
## magnet_belt_x	1.260563	3.42849860	FALSE	FALSE
## magnet_belt_y	1.075472	3.39026255	FALSE	FALSE
## magnet_belt_z	1.102151	4.98343105	FALSE	FALSE
## magnet_arm_x	1.076923	16.25031863	FALSE	FALSE

## magnet_arm_y	1.111111	10.62962019	FALSE	FALSE
## magnet_arm_z	1.024390	15.21794545	FALSE	FALSE
## magnet_dumbbell_x	1.134328	12.09533520	FALSE	FALSE
## magnet_dumbbell_y	1.037975	9.99235279	FALSE	FALSE
## magnet_dumbbell_z	1.223881	8.00407851	FALSE	FALSE
## magnet_forearm_x	1.062500	16.90033138	FALSE	FALSE
## magnet_forearm_y	1.411765	21.34845781	FALSE	FALSE
## magnet_forearm_z	1.035714	18.56997196	FALSE	FALSE
## roll_belt	1.195906	10.33647719	FALSE	FALSE
## roll_arm	45.793103	25.74560285	FALSE	FALSE
## roll_dumbbell	1.096154	89.05174611	FALSE	FALSE
## roll_forearm	10.192053	18.09839409	FALSE	FALSE
## pitch_belt	1.048780	18.36604639	FALSE	FALSE
## pitch_arm	73.833333	28.80448636	FALSE	FALSE
## pitch_dumbbell	2.070175	87.45857762	FALSE	FALSE
## pitch_forearm	61.520000	26.82895743	FALSE	FALSE
## yaw_belt	1.004673	18.65918940	FALSE	FALSE
## yaw_arm	29.511111	28.42212592	FALSE	FALSE
## yaw_dumbbell	1.096154	88.45271476	FALSE	FALSE
## yaw_forearm	14.647619	18.02192200	FALSE	FALSE
## total_accel_belt	1.091532	0.34412439	FALSE	FALSE
## total_accel_arm	1.082386	0.82844762	FALSE	FALSE
## total_accel_dumbbell	1.054250	0.52255927	FALSE	FALSE
## total_accel_forearm	1.094456	0.80295692	FALSE	FALSE
## classe	1.470356	0.06372674	FALSE	FALSE

Finally, it was checked if NAs remained in the data sets :

```
sum(is.na(train2_pred)); sum(is.na(test2_pred))
```

```
## [1] 0
```

```
## [1] 0
```

As both values were zeros, it was possible to conclude that no NAs remained in the data sets.

4. USING RANDOM FORRESTS

4.1. Using Caret package

The prediction model which was first tried used the Caret package and the **random forrests** method. However, due to the size of the `train2_pred` data frame (almost 12,000 entries), the processing time reached several minutes. According to the original paper, using random forrests, a weighted accuracy of 98.2% was achieved.

For this reason, an alternative method for running random forrests was tried, aiming to achieve a satisfactory accuracy.

4.2. Using 'bigrf' package

After a fast research at Google, the **bigrf** package was reverted as a faster option to traditional random forest in Caret package. the **bigrf** package is an implementation of Leo Breiman's and Adele Cutler's Random Forest algorithms for classification and regression, with optimizations for performance and for handling of data sets that are too large to be processed in memory.

- Package source : <https://github.com/aloysius-lim/bigrf>

Additionally, parallel processing using multicore features in `doParallel` package helped to enhance the overall computation speed.

The **bigrf** package could build the prediction model (using random forests method) and classification in few seconds.

4.3. Building the prediction model and applying to training set

The first step was the installation of the new package :

```
if("bigrf" %in% rownames(installed.packages()) == FALSE) {install.packages("bigrf")}
library(bigrf)

## Loading required package: bigmemory
## Loading required package: bigmemory.sri
## Loading required package: BH

## Warning: package 'BH' was built under R version 3.1.2

##
## bigmemory >= 4.0 is a major revision since 3.1.2; please see packages
## biganalytics and and bigtabulate and http://www.bigmemory.org for more information.
```

Later, the parallel/multicore processing was activated :

```
if ("doParallel" %in% rownames(installed.packages()) == FALSE) {
  install.packages("doParallel")
}
library(doParallel)
registerDoParallel(cores = detectCores(all.tests = TRUE))
```

So, the function `bigrfc` was used to build the classification model based in the random forests method :

```
set.seed(1)
fit<-bigrfc(train2_pred, train2_pred$classe,varselect = 1:52)
```

```
## OOB errors:
## Tree Overall error Error by class
##           A           B           C           D           E
## 10           5.96    3.97    8.16    7.74    7.20    3.93
## 20           2.28    0.657  3.774  2.532  3.679  1.709
## 30           1.44    0.329  2.326  1.753  2.591  0.924
## 40           1.16    0.209  1.843  1.509  2.228  0.647
## 50           0.968    0.119  1.667  1.120  2.021  0.462
```

Fifty (50) random forests were performed, in order to reduce the Out-Of-Bag (OOB) classification error to a minimum value.

The function `predict` was used to predict the classes of the same training set used to build the model :

```
pred<-predict(fit, train2_pred, train2_pred$classe)
```

```
## Test errors:
## Tree Overall error Error by class
##           A           B           C           D           E
## 10           0.0255  0.000  0.000  0.000  0.155  0.000
## 20           0.00    0.00  0.00  0.00  0.00  0.00
## 30           0.00    0.00  0.00  0.00  0.00  0.00
## 40           0.00    0.00  0.00  0.00  0.00  0.00
## 50           0.00    0.00  0.00  0.00  0.00  0.00
```

The error rates and confusion matrix can be verified :

```
summary(pred)
```

```
## Predictions on 11776 examples using random forest with 50 trees.
##
## Test set labels:
##
##      A      B      C      D      E
## 3348 2279 2054 1930 2165
##
## Overall error rate: 0.00
##
## Test set confusion matrix (OOB):
##      Predicted
## Actual      A      B      C      D      E
##      A 3348      0      0      0      0
##      B      0 2279      0      0      0
##      C      0      0 2054      0      0
##      D      0      0      0 1930      0
##      E      0      0      0      0 2165
```

It was possible to verify that the built model scored **100%** accuracy of prediction in all classes, and **0.00%** overall error, in the training data set.

4.4. Testing prediction model using cross-validation

In spite of the built model showing good results in the prediction of training data set, an additional and important test was to apply the same model to a new data set. For this purpose, the `test2_pred` data set was used :

```
pred2<-predict(fit, test2_pred, test2_pred$classe)
```

```
## Test errors:
## Tree Overall error Error by class
##      A      B      C      D      E
## 10      1.50 0.986 2.503 1.974 1.944 0.416
## 20      0.905 0.314 1.713 1.023 1.477 0.347
## 30      0.803 0.224 1.120 0.877 1.866 0.347
## 40      0.701 0.224 0.922 0.804 1.477 0.416
## 50      0.688 0.224 0.856 0.731 1.555 0.416
```

Finally, the error rates can be verified and a confusion matrix could be built, to compare the actual values to its predicted values :

```
summary(pred2)
```

```
## Predictions on 7846 examples using random forest with 50 trees.
##
## Test set labels:
##
##      A      B      C      D      E
## 2232 1518 1368 1286 1442
##
## Overall error rate: 0.688
##
## Test set confusion matrix (OOB):
##      Predicted
```

## Actual	A	B	C	D	E
## A	2227	3	1	0	1
## B	6	1505	7	0	0
## C	0	8	1358	2	0
## D	0	0	20	1266	0
## E	0	0	0	6	1436

The out of sample error rate in the classification was **0.726%**, which can be considered low.

The accuracy of the model applied to `test2_pred` was :

- Class A : 99.8%
- Class B : 99.1%
- Class C : 99.3%
- Class D : 98.4%
- Class E : 99.6%
- **Overall accuracy : 99.3%**

The overall accuracy of **99.3%** was considered to be suitable, and for this reason, the built model was accepted to the next step.

4.5. Applying the prediction model to test set

Finally, the built prediction model was applied to the `test` data set.

```
test_pred<-test[, preds]
pred3<-predict(fit, test_pred)
```

```
## Processing tree number:
## 10
## 20
## 30
## 40
## 50
```

```
pred3<-replace(pred3, pred3==1, "A")
pred3<-replace(pred3, pred3==2, "B")
pred3<-replace(pred3, pred3==3, "C")
pred3<-replace(pred3, pred3==4, "D")
pred3<-replace(pred3, pred3==5, "E")
```

The predicted classes were :

```
as.data.frame(pred3)
```

```
## pred3
## 1 B
## 2 A
## 3 B
## 4 A
## 5 A
## 6 E
## 7 D
```

## 8	B
## 9	A
## 10	A
## 11	B
## 12	C
## 13	B
## 14	A
## 15	E
## 16	E
## 17	A
## 18	B
## 19	B
## 20	B

5. CONCLUSION

The random forests method is, indeed, very accurate in classification procedures. The computing time may become an issue when the data set is large. Fortunately, there are options to make this process faster, keeping its accuracy in a high level.

The predicted classes for `test` data set were all correct according to project submission system, which confirms the high accuracy of the proposed process.

It was possible to achieve a better accuracy than obtained in the original paper. As seen in item 4.4, the overall accuracy was approximately **99.3%**, against **98.2%** of the original paper.

The out of sample error was low, only **0.726%**.

Finally, the parallel computation using multicore features of modern processors is a very relevant procedure for larger scale computations.